



## AGENTIC AI LABORATORY

**Presented by:**  
**Yashwi Pandey**

**Under the supervision of**  
**Mr. Ayush Kumar Singh**

**BACHELOR OF TECHNOLOGY**

**In**

**COMPUTER SCIENCE AND ENGINEERING**

**SHARDA SCHOOL OF ENGINEERING AND TECHNOLOGY**

**GREATER NOIDA**

# Fine-Tuning a Small Language Model (SLM) Using LoRA and 4-bit Quantization

## 1. Introduction

This project demonstrates how to fine-tune a Small Language Model (SLM) efficiently using modern parameter-efficient fine-tuning techniques. The base model used is TinyLlama-1.1B-Chat, and it is fine-tuned on a dataset of English quotes.

To reduce memory usage and computational cost, the project applies:

- 4-bit quantization (QLoRA)
- LoRA (Low-Rank Adaptation)
- Supervised Fine-Tuning (SFT) using the TRL library

The final model is evaluated using both qualitative text generation and quantitative metrics such as loss and perplexity.

## 2. Library Installation

This cell installs all required libraries:

- PyTorch – Core deep learning framework
- Transformers – Pretrained models and tokenizers
- Datasets – Loading and processing datasets
- PEFT – Parameter-Efficient Fine-Tuning (LoRA)
- BitsAndBytes – 4-bit and 8-bit quantization
- TRL – Training utilities for language models

These libraries together enable efficient fine-tuning on limited hardware.

## 3. Imports and Configuration:

```
import torch
from datasets import load_dataset
from transformers import AutoModelForCausalLM, AutoTokenizer, BitsAndBytesConfig
from peft import LoraConfig
from trl import SFTTrainer
```

Model and Dataset Selection:

```
model_name = "TinyLlama/TinyLlama-1.1B-Chat-v1.0"
dataset_name = "Abirate/english_quotes"
new_model = "TinyLlama-1.1B-Quotes-Finetuned"
```

**Base Model:** TinyLlama (1.1 billion parameters)

**Dataset:** English quotes dataset

**Output Model:** Name used to save the fine-tuned model

## **4. Dataset Loading**

```
dataset = load_dataset(dataset_name, split="train")
```

The dataset is loaded using the Hugging Face datasets library.

Each training sample contains a quote, which is used as the text for supervised fine-tuning.

## **5. 4-bit Quantization Configuration**

### **Purpose of Quantization**

- Reduces memory usage dramatically
- Enables training large models on consumer GPUs

### **Key Settings**

- NF4: Normal Float 4-bit quantization
- float16 computation: Improves speed and stability
- QLoRA approach: Quantized base model + trainable LoRA adapters

## **6. Loading the Base Model**

```
model = AutoModelForCausalLM.from_pretrained(  
    model_name,  
    quantization_config=bnb_config,  
    device_map={"": 0}  
)
```

Loads the pretrained TinyLlama model

Applies 4-bit quantization

Automatically places the model on GPU

Additional configuration:

```
model.config.use_cache = False
```

```
model.config.pretraining_tp = 1  
Disabling cache avoids training issues
```

Ensures compatibility with LoRA fine-tuning

## 7. Tokenizer Setup

- Uses the same tokenizer as the base model
- Sets padding token to EOS token (required for causal LM training)
- Right padding is optimal for transformer models

## 8. LoRA (Low-Rank Adaptation) Configuration

### Why LoRA?

- Only a small number of parameters are trained
- Base model weights remain frozen
- Faster training and lower memory usage

### Key Parameters

- **r = 64:** Rank of LoRA matrices
- **alpha = 16:** Scaling factor
- **dropout = 0.1:** Regularization

## 9. Supervised Fine-Tuning (SFT) Configuration

### Training Strategy

- **SFT (Supervised Fine-Tuning):** Model learns directly from labeled text
- **Single epoch:** Sufficient for demonstration purposes
- **Batch size of 4:** Balanced for memory and speed

Additional settings control:

- Logging frequency
- Optimizer ([paged\\_adamw\\_32bit](#))
- Gradient stability
- Learning rate scheduling

## 10. Trainer Initialization and Training

The [SFTTrainer](#):

- Handles tokenization
- Applies LoRA adapters
- Manages training loop

## 11. Saving the Fine-Tuned Model

The fine-tuned LoRA-adapted model is saved locally for future inference or deployment.

## 12. Qualitative Evaluation (Inference)

The model is tested using a natural language prompt to evaluate:

- Coherence
- Relevance
- Fluency of generated quotes

This provides a human-interpretable evaluation of training quality.

## 13. Quantitative Evaluation: Loss and Perplexity

### Loss Extraction

```
python  
  
log_history = trainer.state.log_history
```

- Training logs are parsed to extract loss values
- Final loss represents model convergence

### Perplexity Calculation

```
python  
  
perplexity = exp(final_loss)
```

- Lower perplexity indicates better language modeling
- Measures how “surprised” the model is by the text

## 14. Training Loss Visualization

A loss curve is plotted over training steps  
Helps verify stable learning and convergence

## **15. Final Text Generation Test**

The final generation test confirms:

- The model has learned stylistic patterns of quotes
- The output is meaningful and contextually appropriate

## **16. Conclusion**

This project successfully demonstrates:

- Fine-tuning a Small Language Model using LoRA
- Efficient training with 4-bit quantization
- Practical evaluation using loss, perplexity, and text generation

**The approach is scalable, memory-efficient, and suitable for real-world NLP applications where computational resources are limited.**